

Principal Component Analysis (PCA)

Introduction

In machine learning and data analytics, datasets often contain a large number of features. Some of these features may be highly related and provide similar information. Working with many features increases computational complexity and may affect the efficiency of machine learning models.

Principal Component Analysis (PCA) is a dimensionality reduction technique that reduces the number of features in a dataset while retaining most of the important information. PCA transforms the original features into a smaller set of new features called **Principal Components**.

Dimensionality Reduction

Dimensionality reduction is the process of reducing the number of input variables or features in a dataset while preserving as much useful information as possible.

For example, consider a student dataset containing marks in Mathematics, Physics, Chemistry, Biology, and Computer Science. Since some subjects may be closely related, the information contained in multiple features can often be represented using fewer features. Dimensionality reduction helps achieve this objective.

Definition of PCA

Principal Component Analysis (PCA) is an unsupervised learning technique used to transform a high-dimensional dataset into a lower-dimensional dataset. It identifies patterns in data and finds directions along which the variation in the data is maximum.

The transformed variables are known as Principal Components. These components are linear combinations of the original features and are arranged in descending order of importance.

Understanding PCA Through an Example

Consider a dataset containing marks in Mathematics, Physics, and Chemistry. Students who score highly in Mathematics often perform well in Physics and Chemistry. Since these features are correlated, they may contain overlapping information.

Instead of using all three features separately, PCA combines the information and creates new variables called Principal Components.

Original Features

- Mathematics
- Physics

- Chemistry

Transformed Features

- Principal Component 1 (PC1)
- Principal Component 2 (PC2)

Thus, PCA reduces the number of features while preserving most of the information present in the dataset.

Principal Components

Principal Components are the new variables generated by PCA.

Principal Component 1 (PC1)

PC1 captures the maximum variation present in the dataset. It contains the most significant information and explains the largest portion of the data variability.

Principal Component 2 (PC2)

PC2 captures the second highest variation in the dataset and is independent of PC1.

Principal Component 3 (PC3)

PC3 captures the remaining variation and is independent of both PC1 and PC2.

The importance of principal components decreases from PC1 to PC2 to PC3 and so on.

Need for PCA

PCA is widely used because it helps to:

- Reduce the number of features
 - Eliminate redundant information
 - Improve computational efficiency
 - Reduce storage requirements
 - Simplify data visualization
 - Improve model performance
 - Reduce overfitting in machine learning models
-

Working of PCA

The PCA process involves the following steps:

Step 1

Collect and prepare the dataset.

Step 2

Standardize the dataset so that all features have the same scale.

Step 3

Calculate the covariance matrix to determine the relationships among features.

Step 4

Compute Eigenvalues and Eigenvectors.

Step 5

Select the principal components having the highest Eigenvalues.

Step 6

Transform the original data into a new lower-dimensional feature space.

Standardization in PCA

Before applying PCA, data should be standardized because different features may have different ranges.

For example:

Feature	Range
Age	18–60
Salary	20,000–100,000

Without standardization, features with larger values may dominate the PCA calculation.

Standardization converts all features to a common scale with:

- Mean equal to 0
 - Standard deviation equal to 1
-

Eigenvalues and Eigenvectors

Eigenvectors

Eigenvectors determine the directions of the principal components.

Eigenvalues

Eigenvalues indicate the amount of information captured by each principal component.

A higher Eigenvalue indicates that the corresponding principal component contains more useful information.

Explained Variance Ratio

The Explained Variance Ratio indicates how much information is retained by each principal component.

Example:

Principal Component	Variance Explained
PC1	70%
PC2	20%
PC3	10%

Using only PC1 and PC2 retains:

$$70\% + 20\% = 90\%$$

This means that two principal components preserve 90% of the information present in the original dataset.

Feature Selection and PCA

Feature Selection	PCA
Selects existing features	Creates new features
Original features remain unchanged	Original features are transformed
Easy to interpret	Difficult to interpret
Reduces features by selection	Reduces features by transformation

Advantages of PCA

- Reduces dimensionality
 - Removes redundancy among features
 - Improves computational speed
 - Reduces storage requirements
 - Simplifies data visualization
 - Helps in handling high-dimensional datasets
 - Reduces overfitting
-

Limitations of PCA

- Principal components are difficult to interpret
 - Some information may be lost
 - Suitable only for numerical data
 - Sensitive to feature scaling
 - May not capture complex nonlinear relationships
-

Applications of PCA

Image Processing

Used for image compression and facial recognition systems.

Healthcare

Used in disease prediction and medical image analysis.

Finance

Used in stock market analysis and risk management.

Marketing

Used for customer segmentation and recommendation systems.

Data Visualization

Used to visualize high-dimensional datasets in two-dimensional or three-dimensional space.

PCA Using Python

```
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import pandas as pd
```

```
df = pd.read_csv("data.csv")

scaler = StandardScaler()
X_scaled = scaler.fit_transform(df)

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

print(X_pca)
```

Summary

Principal Component Analysis is one of the most widely used dimensionality reduction techniques in machine learning. It transforms a large set of correlated features into a smaller set of uncorrelated principal components while preserving most of the important information. PCA improves computational efficiency, reduces storage requirements, simplifies visualization, and enhances the performance of machine learning models. As a result, PCA is an important preprocessing technique used in various domains such as healthcare, finance, image processing, and data analytics.